

Static Java, GraalVM Native and OpenJDK

Andrew Dinn
Red Hat Java Team
October 2021

Static Java?

Static Java?

- **A New Deployment Model?**
 - Pre-compiled, self-contained executable like C++/Go

Static Java?

- **A New Deployment Model?**
 - Pre-compiled, self-contained executable like C++/Go
 - No class files → no interpreter/JIT, less metadata ...

Static Java?

- **A New Deployment Model?**
 - Pre-compiled, self-contained executable like C++/Go
 - No class files → no interpreter/JIT, less metadata ...
- **A Container-Friendly Java?**
 - Smaller footprint/faster startup

Static Java?

- **A New Deployment Model?**
 - Pre-compiled, self-contained executable like C++/Go
 - No class files → no interpreter/JIT, less metadata ...
- **A Container-Friendly Java?**
 - Smaller footprint/faster startup
 - Not automatic

Static Java?

- **A New Deployment Model?**
 - Pre-compiled, self-contained executable like C++/Go
 - No class files → no interpreter/JIT, less metadata ...
- **A Container-Friendly Java?**
 - Smaller footprint/faster startup
 - Not automatic
- **A Reduced Java?**
 - **Dynamic** (runtime) → **Static** (buildtime)

Static Java?

- **A New Deployment Model?**
 - Pre-compiled, self-contained executable like C++/Go
 - No class files → no interpreter/JIT, less metadata ...
- **A Container-Friendly Java?**
 - Smaller footprint/faster startup
 - Not automatic
- **A Reduced Java?**
 - Dynamic (runtime) → Static (buildtime)
 - Java? or just the JVM/JDK?

Dynamic Java

How Can Java Be Static?

- Program **code** evolves during execution

How Can Java Be Static?

- **Program code evolves during execution**
 - Class loading, Service providers, Class generators

How Can Java Be Static?

- **Program code evolves during execution**
 - Class loading, Service providers, Class generators
 - Loader Delegation, Module Layers

How Can Java Be Static?

- **Program code evolves during execution**
 - Class loading, Service providers, Class generators
 - Loader Delegation, Module Layers
- **Execution model supports this evolution**

How Can Java Be Static?

- **Program code evolves during execution**
 - Class loading, Service providers, Class generators
 - Loader Delegation, Module Layers
- **Execution model supports this evolution**
 - Reflection, Method/VarHandles

How Can Java Be Static?

- **Program code evolves during execution**
 - Class loading, Service providers, Class generators
 - Loader Delegation, Module Layers
- **Execution model supports this evolution**
 - Reflection, Method/VarHandles
 - With ClassLoader → indirect load and link

How Can Java Be Static?

- Program **state** evolves too

How Can Java Be Static?

- Program state evolves too
 - Load → Static Init → Interpret → (Re-)Compile

How Can Java Be Static?

- Program state evolves too
 - { Load → Static Init → Interpret → (Re-)Compile } +

How Can Java Be Static?

- Program state evolves too
 - { Load → Static Init → Interpret → (Re-)Compile } +
 - n.b. Compile happens after Static Init/Execute

How Can Java Be Static?

- Program state evolves too
 - { Load → Static Init → Interpret → (Re-)Compile } +
 - n.b. Compile happens after Static Init/Execute
- Static Init may recursively embed this pattern

How Can Java Be Static?

- Program state evolves too
 - { Load → Static Init → Interpret → (Re-)Compile } +
 - n.b. Compile happens after Static Init/Execute
- Static Init may recursively embed this pattern
 - Static Init → { Load → ... } +

How Can Java Be Static?

- Program state evolves too
 - { Load → Static Init → Interpret → (Re-)Compile } +
 - n.b. Compile happens after Static Init/Execute
- Static Init may recursively embed this pattern
 - Static Init → { Load → ... } +
- Load also drives recompile

How Can Java Be Static?

- Program state evolves too
 - { Load → Static Init → Interpret → (Re-)Compile } +
 - n.b. Compile happens after Static Init/Execute
- Static Init may recursively embed this pattern
 - Static Init → { Load → ... } +
- Load also drives recompile
 - Load → Deopt(dependent) → Interpret → ReCompile

How Can Java Be Static?

- VM configured at runtime
 - GC, JITs, Compressed Oops, HW Accelerators, etc

How Can Java Be Static?

- **VM configured at runtime**
 - GC, JITs, Compressed Oops, HW Accelerators, etc
 - Configurable on command line

How Can Java Be Static?

- **VM configured at runtime**
 - GC, JITs, Compressed Oops, HW Accelerators, etc
 - Configurable on command line
 - Defaulted according to runtime env

How Can Java Be Static?

- **VM configured at runtime**
 - GC, JITs, Compressed Oops, HW Accelerators, etc
 - Configurable on command line
 - Defaulted according to runtime env
- **Tools**
 - Agents, JDWP, JFR, JMX, jcmd, etc

How Can Java Be Static?

- **VM configured at runtime**
 - GC, JITs, Compressed Oops, HW Accelerators, etc
 - Configurable on command line
 - Defaulted according to runtime env
- **Tools**
 - Agents, JDWP, JFR, JMX, jcmd, etc
 - Can be installed & configured at runtime

How Can Java Be Static?

- **VM configured at runtime**
 - GC, JITs, Compressed Oops, HW Accelerators, etc
 - Configurable on command line
 - Defaulted according to runtime env
- **Tools**
 - Agents, JDWP, JFR, JMX, jcmd, etc
 - Can be installed & configured at runtime
 - All rely on dynamic capabilities

How Can We Make Java Static?

How To Make Java Static?

- Close the world of the app
 - All code identified and compiled ahead of time

How To Make Java Static?

- Close the world of the app
 - All code identified and compiled ahead of time
 - Trace class, method and field references from App main

How To Make Java Static?

- **Close the world of the app**
 - All code identified and compiled ahead of time
 - Trace class, method and field references from App main
 - Try not to include the kitchen sink

How To Make Java Static?

- **Close the world of the app**
 - All code identified and compiled ahead of time
 - Trace class, method and field references from App main
 - Try not to include the kitchen sink
- **Configure indirectly accessed code?**

How To Make Java Static?

- **Close the world of the app**
 - All code identified and compiled ahead of time
 - Trace class, method and field references from App main
 - Try not to include the kitchen sink
- **Configure indirectly accessed code?**
 - Reflection & method/varhandle targets

How To Make Java Static?

- **Close the world of the app**
 - All code identified and compiled ahead of time
 - Trace class, method and field references from App main
 - Try not to include the kitchen sink
- **Configure indirectly accessed code?**
 - Reflection & method/varhandle targets
 - Service implementations

How To Make Java Static?

- **Close the world of the app**
 - All code identified and compiled ahead of time
 - Trace class, method and field references from App main
 - Try not to include the kitchen sink
- **Configure indirectly accessed code?**
 - Reflection & method/varhandle targets
 - Service implementations
- **Configure classpath resources?**

How To Make Java Static?

- Close the world of the **target** runtime

How To Make Java Static?

- Close the world of the target runtime
 - But do it at **build** time

How To Make Java Static?

- Close the world of the target runtime
 - But do it at build time
- **Can** we emulate the dynamic runtime?

How To Make Java Static?

- Close the world of the target runtime
 - But do it at build time
- Can we emulate the dynamic runtime?
 - App-defined loaders in build env?

How To Make Java Static?

- Close the world of the target runtime
 - But do it at build time
- Can we emulate the dynamic runtime?
 - App-defined loaders in build env?
 - Or just the classpath & modulepath?

How To Make Java Static?

- Close the world of the target runtime
 - But do it at build time
- Can we emulate the dynamic runtime?
 - App-defined loaders in build env?
 - Or just the classpath & modulepath?
 - Which services get loaded?

How To Make Java Static?

- Close the world of the target runtime
 - But do it at build time
- Can we emulate the dynamic runtime?
 - App-defined loaders in build env?
 - Or just the classpath & modulepath?
 - Which services get loaded?
 - e.g. what will the target LANG setting be

How To Make Java Static?

- No Load → No static init at load?

How To Make Java Static?

- No Load → No static init at load?
 - So when do we run class static init?

How To Make Java Static?

- No Load → No static init at load?

1 Emulate dynamic JVM?

How To Make Java Static?

- No Load → No static init at load?
- 1 Emulate dynamic JVM?
 - Run init at point(s) where load *would* have happened?
 - Including init that precedes App main

How To Make Java Static?

- No Load → No static init at load?

1 Emulate dynamic JVM?

- Run init at point(s) where load *would* have happened?
 - Including init that precedes App main
 - Check at every new or static access/call
 - Or at mention of `Foo.class` or `forName()` call

How To Make Java Static?

- No Load → No static init at load?

1 Emulate dynamic JVM?

- Run init at point(s) where load *would* have happened?
 - Including init that precedes App main
 - Check at every new or static access/call
 - Or at mention of Foo.class or `forName()` call
 - Init only once (however we get there)

How To Make Java Static?

- No Load → No static init at load?

1 Emulate dynamic JVM?

- Run init at point(s) where load *would* have happened?
 - Including init that precedes App main
 - Check at every new or static access/call
 - Or at mention of Foo.class or `forName()` call
 - Init only once (however we get there)
 - Analysis can remove many redundant attempts

How To Make Java Static?

- No Load → No static init at load
- 2 Run all inits in one go before App main?

How To Make Java Static?

- No Load → No static init at load
- 2 Run all inits in one go before App main?
 - But still in a 'correct' order?

How To Make Java Static?

- No Load → No static init at load
- ## 2 Run all inits in one go before App main?
- But still in a 'correct' order?
 - Might this change the init outcome relative to 1?
`A::init(); A.STATIC_VALUE = X; B::init()`

How To Make Java Static?

- No Load → No static init at load
- 2 Run all inits in one go before App main?
- But still in a 'correct' order?
 - Might this change the init outcome relative to 1?
`A::init(); A.STATIC_VALUE = X; B::init()`
 - Maybe some hybrid of 1 and 2 could avoid that

How To Make Java Static?

- No Load → No static init at load
- ## 2 Run all inits in one go before App main?
- But still in a 'correct' order?
 - Might this change the init outcome relative to 1?
`A::init(); A.STATIC_VALUE = X; B::init()`
 - Maybe some hybrid of 1 and 2 could avoid that
 - Needs a very careful analysis

How To Make Java Static?

- No Load → No static init at load
- 3 Pre-compute static fields at build time?

How To Make Java Static?

- No Load → No static init at load
- ## 3 Pre-compute static fields at build time?
- As appropriate to the **target** env

How To Make Java Static?

- No Load → No static init at load
- ## 3 Pre-compute static fields at build time?
- As appropriate to the target env
 - See above ...

How To Make Java Static?

- No Load → No static init at load
- ## 3 Pre-compute static fields at build time?
- As appropriate to the target env
 - See above ...
... but with clairvoyance requirement

How To Make Java Static?

- No Load → No static init at load
- ### 3 Pre-compute static fields at build time?
- As appropriate to the target env
 - See above ...
... but with clairvoyance requirement
 - Actually a hybrid approach **must** be adopted ...

How To Make Java Static?

- No Load → No static init at load
- ## 3 Pre-compute static fields at build time?
- As appropriate to the target env
 - See above ...
... but with clairvoyance requirement
 - Actually a hybrid approach must be adopted ...
 - ... and sometimes you need to cheat?

How To Make Java Static?

- Build time init is **fine** for many fields

How To Make Java Static?

- Build time init is fine for many fields
 - Final primitive constants

How To Make Java Static?

- **Build time init is fine for many fields**
 - Final primitive constants
 - Immutable objects (e.g. Strings)

How To Make Java Static?

- Build time init is fine for many fields
 - Final primitive constants
 - Immutable objects (e.g. Strings)
 - Final mutable objects?

How To Make Java Static?

- Build time init is fine for many fields
 - Final primitive constants
 - Immutable objects (e.g. Strings)
 - Final mutable objects?
- BTI even opens up new compiler optimizations

How To Make Java Static?

- **Build time init is fine for many fields**
 - Final primitive constants
 - Immutable objects (e.g. Strings)
 - Final mutable objects?
- **BTI even opens up new compiler optimizations**
 - Non-final values that are read-only in app closure
 - Effectively final

How To Make Java Static?

- But some values really must be **runtime** initied

How To Make Java Static?

- But some values really must be runtime initied
 - Thread, FileStream

How To Make Java Static?

- But some values really must be runtime initied
 - Thread, FileStream
 - Random, Credentials, etc

How To Make Java Static?

- But some values really must be runtime initied
 - Thread, FileStream
 - Random, Credentials, etc
 - Values computed from native calls/JVM/OS env

How To Make Java Static?

- But some values really must be runtime initied
 - Thread, FileStream
 - Random, Credentials, etc
 - Values computed from native calls/JVM/OS env
 - Values derived from any/all of the above

How To Make Java Static?

- It matters **how** values are computed/consumed

How To Make Java Static?

- It matters how values are computed/consumed
 - `public static final int
UnsafeConstants.PAGE_SIZE = 0;`

How To Make Java Static?

- It matters how values are computed/consumed
 - `public static final int
UnsafeConstants.PAGE_SIZE = 0;`
 - JVM bootstrap injects real value

How To Make Java Static?

- It matters how values are computed/consumed
 - `public static final int UnsafeConstants.PAGE_SIZE = 0;`
 - JVM bootstrap injects real value
 - Used/exposed via other APIs? e.g. ByteBuffer?

How To Make Java Static?

- It matters how values are computed/consumed
 - `public static final int UnsafeConstants.PAGE_SIZE = 0;`
 - JVM bootstrap injects real value
 - Used/exposed via other APIs? e.g. `ByteBuffer`?
 - Default/Alternative Locales

How To Make Java Static?

- It matters how values are computed/consumed
 - `public static final int UnsafeConstants.PAGE_SIZE = 0;`
 - JVM bootstrap injects real value
 - Used/exposed via other APIs? e.g. `ByteBuffer`?
 - Default/Alternative Locales
 - Choice drives init of text processing subsystem

How To Make Java Static?

- It matters how values are computed/consumed
 - `public static final int UnsafeConstants.PAGE_SIZE = 0;`
 - JVM bootstrap injects real value
 - Used/exposed via other APIs? e.g. `ByteBuffer`?
 - Default/Alternative Locales
 - Choice drives init of text processing subsystem
 - Default `FileSystemProvider`

How To Make Java Static?

- It matters how values are computed/consumed
 - `public static final int UnsafeConstants.PAGE_SIZE = 0;`
 - JVM bootstrap injects real value
 - Used/exposed via other APIs? e.g. `ByteBuffer`?
 - Default/Alternative Locales
 - Choice drives init of text processing subsystem
 - Default `FileSystemProvider`
 - Associated `FileSystem` caches **local** root/working dir

How To Make Java Static?

- **JVM Configuration**
 - **Configure compilation for target JVM environment**

How To Make Java Static?

- **JVM Configuration**
 - **Configure compilation for target JVM environment**
 - **GC, Compressed Oops, HW Accelerators, etc**

How To Make Java Static?

- **JVM Configuration**
 - **Configure compilation for target JVM environment**
 - GC, Compressed Oops, HW Accelerators, etc
- **Tools Configuration**
 - **Convert some tools to operate at build time?**

How To Make Java Static?

- **JVM Configuration**
 - Configure compilation for target JVM environment
 - GC, Compressed Oops, HW Accelerators, etc
- **Tools Configuration**
 - Convert some tools to operate at build time?
 - Agents, JFR

How To Make Java Static?

- **JVM Configuration**
 - Configure compilation for target JVM environment
 - GC, Compressed Oops, HW Accelerators, etc
- **Tools Configuration**
 - Convert some tools to operate at build time?
 - Agents, JFR
 - Graal: JFR with a runtime on/off switch

How To Make Java Static?

- **JVM Configuration**
 - Configure compilation for target JVM environment
 - GC, Compressed Oops, HW Accelerators, etc
- **Tools Configuration**
 - Convert some tools to operate at build time?
 - Agents, JFR
 - Graal: JFR with a runtime on/off switch
 - Implement complete replacements for others?

How To Make Java Static?

- **JVM Configuration**
 - Configure compilation for target JVM environment
 - GC, Compressed Oops, HW Accelerators, etc
- **Tools Configuration**
 - Convert some tools to operate at build time?
 - Agents, JFR
 - Graal: JFR with a runtime on/off switch
 - Implement complete replacements for others?
 - Graal: JDWP → DWARF, JMX/jcmd → ???

GraalVM Native

GraalVM Native

- Closure Analysis
 - Derives full reference closure

GraalVM Native

- **Closure Analysis**
 - Derives full reference closure
 - Able to prune unused code and (instance/static) state
 - More or less aggressively (high build time cost)

GraalVM Native

- **Closure Analysis**
 - Derives full reference closure
 - Able to prune unused code and (instance/static) state
 - More or less aggressively (high build time cost)
- **Class/Module Loading**
 - Standard ClassLoaders exist but don't load

GraalVM Native

- **Closure Analysis**
 - Derives full reference closure
 - Able to prune unused code and (instance/static) state
 - More or less aggressively (high build time cost)
- **Class/Module Loading**
 - Standard ClassLoaders exist but don't load
 - Return existing class or throw CNFE

GraalVM Native

- **Closure Analysis**
 - Derives full reference closure
 - Able to prune unused code and (instance/static) state
 - More or less aggressively (high build time cost)
- **Class/Module Loading**
 - Standard ClassLoaders exist but don't load
 - Return existing class or throw CNFE
 - Queryable Module base layer info recently added

GraalVM Native

- **Reflection**
 - Analysis infers Method/Field target where possible

GraalVM Native

- **Reflection**
 - Analysis infers Method/Field target where possible
 - Fall back to user configuration

GraalVM Native

- **Reflection**
 - Analysis infers Method/Field target where possible
 - Fall back to user configuration
- **Method Handles**
 - Handle chains must end at `DirectHandle`

GraalVM Native

- **Reflection**
 - Analysis infers Method/Field target where possible
 - Fall back to user configuration
- **Method Handles**
 - Handle chains must end at `DirectHandle`
 - Lambdas still work

GraalVM Native

- **Reflection**
 - Analysis infers Method/Field target where possible
 - Fall back to user configuration
- **Method Handles**
 - Handle chains must end at `DirectHandle`
 - Lambdas still work
- **Service & Resource Loading**
 - Analysis infers target where possible

GraalVM Native

- **Reflection**
 - Analysis infers Method/Field target where possible
 - Fall back to user configuration
- **Method Handles**
 - Handle chains must end at `DirectHandle`
 - Lambdas still work
- **Service & Resource Loading**
 - Analysis infers target where possible
 - Fall back to user configuration

GraalVM Native

- **Special cased JDK Service Providers**
 - **Locale Support**

GraalVM Native

- **Special cased JDK Service Providers**
 - **Locale Support**
 - **User configures default locale at build**

GraalVM Native

- Special cased JDK Service Providers
 - Locale Support
 - User configures default locale at build
 - Must also pre-specify **all** other required locales

GraalVM Native

- Special cased JDK Service Providers
 - Locale Support
 - User configures default locale at build
 - Must also pre-specify **all** other required locales
 - Default File System

GraalVM Native

- **Special cased JDK Service Providers**
 - **Locale Support**
 - User configures default locale at build
 - Must also pre-specify **all** other required locales
 - **Default File System**
 - Bootstrap requires default Provider & FileSystem

GraalVM Native

- **Special cased JDK Service Providers**
 - **Locale Support**
 - User configures default locale at build
 - Must also pre-specify **all** other required locales
 - **Default File System**
 - Bootstrap requires default Provider & FileSystem
 - “Optimized” by creating at build-time

GraalVM Native

- Special cased JDK Service Providers
 - Locale Support
 - User configures default locale at build
 - Must also pre-specify **all** other required locales
 - Default File System
 - Bootstrap requires default Provider & FileSystem
 - “Optimized” by creating at build-time
 - Requires **patching** JDK class

GraalVM Native

- Almost all of JDK is build-time init'd

GraalVM Native

- Almost all of JDK is build-time init'd
 - Some of it is always runtime init'd

GraalVM Native

- Almost all of JDK is build-time init'd
 - Some of it is always runtime init'd
 - Some of it gets init'd twice

GraalVM Native

- Almost all of JDK is build-time init'd
 - Some of it is always runtime init'd
 - Some of it gets init'd twice
 - Which is a certainly questionable?

GraalVM Native

- Almost all of JDK is build-time init'd
 - Some of it is always runtime init'd
 - Some of it gets init'd twice
 - Which is a certainly questionable?
- Apps **can** be configured for BTI

GraalVM Native

- Almost all of JDK is build-time init'd
 - Some of it is always runtime init'd
 - Some of it gets init'd twice
 - Which is a certainly questionable?
- Apps can be configured for BTI
 - Almost **all** frameworks try to do this

GraalVM Native

- **Almost all of JDK is build-time init**
 - Some of it is always runtime init
 - Some of it gets init twice
 - Which is a certainly questionable?
- **Apps can be configured for BTI**
 - Almost all frameworks try to do this
- **Static init analysis rejects invalid cases**

GraalVM Native

- Almost all of JDK is build-time init'd
 - Some of it is always runtime init'd
 - Some of it gets re-init'd
 - Which is a certainly questionable?
- Apps can be configured for BTI
 - Almost all frameworks try to do this
- Static init analysis rejects invalid cases
 - Values that **must** be RTI

GraalVM Native

- Almost all of JDK is build-time init'd
 - Some of it is always runtime init'd
 - Some of it gets re-init'd
 - Which is a certainly questionable?
- Apps can be configured for BTI
 - Almost all frameworks try to do this
- Static init analysis rejects invalid cases
 - Values that must be RTI
 - Values that **derive** from RTI values

GraalVM Native

- BTI is critical to small JDK footprint/fast startup

GraalVM Native

- **BTI is critical to small JDK footprint/fast startup**
 - Only static state referenced by app is retained

GraalVM Native

- **BTI is critical to small JDK footprint/fast startup**
 - Only static state referenced by app is retained
 - Intermediate computed values squeezed out
 - As is their type info

GraalVM Native

- **BTI is critical to small JDK footprint/fast startup**
 - Only static state referenced by app is retained
 - Intermediate computed values squeezed out
 - As is their type info
 - Only RTI static init code is compiled/executed

GraalVM Native

- **BTI is critical to small JDK footprint/fast startup**
 - Only static state referenced by app is retained
 - Intermediate computed values squeezed out
 - As is their type info
 - Only RTI static init code is compiled/executed
 - Especially JDK default locale and file system setup

GraalVM Native

- **BTI is critical to small JDK footprint/fast startup**
 - Only static state referenced by app is retained
 - Intermediate computed values squeezed out
 - As is their type info
 - Only RTI static init code is compiled/executed
 - Especially JDK default locale and file system setup
 - This is a big slice of dynamic JVM startup

GraalVM Native

- **BTI is critical to small JDK footprint/fast startup**
 - Only static state referenced by app is retained
 - Intermediate computed values squeezed out
 - As is their type info
 - Only RTI static init code is compiled/executed
 - Especially JDK default locale and file system setup
 - This is a big slice of dynamic JVM startup
 - n.b. the same story applies for app code

GraalVM Native

- Example of REINIT

```
class JDKInitializationFeature {  
    ...  
    RuntimeClassInitializationSupport rci = ...;  
    rci.initializeAtBuildTime("java.io",  
                             "JDK core init'd at build time");  
    rci.initializeAtBuildTime("java.lang",  
                             "JDK core init'd at build time");  
    ...  
    rci.rerunInitialization("sun.nio.ch.Epoll",  
                           "Calls Epoll.eventSize(), ...");  
}
```

GraalVM Native

- Example of REINIT

```
class Epoll {  
    ...  
    private static final  
    int SIZEOF_EPOLLEVENT    = eventSize();  
    ...  
    // opcodes  
    static final int EPOLL_CTL_ADD = 1;  
    ...  
    private static native int eventSize();  
}
```

GraalVM Native

- Example of REINIT

```
class Epoll {  
    ...  
    private static final  
    int SIZEOF_EPOLLEVENT    = eventSize();  
    ...  
    // opcodes  
    static final int EPOLL_CTL_ADD = 1;  
    ...  
    private static native int eventSize();  
}
```

GraalVM Native

- Example of REINIT

```
class Epoll {  
    ...  
    private static final  
    int SIZEOF_EPOLLEVENT    = eventSize();  
    ...  
    // opcodes  
    static final int EPOLL_CTL_ADD = 1;  
    ...  
    private static native int eventSize();
```

GraalVM Native

- BTI is occasionally **desirable** and **wrong**

GraalVM Native

- BTI is occasionally **desirable** and **wrong**
 - BTI of `java.base` requires a local `FileSystem` instance

GraalVM Native

- BTI is occasionally **desirable** and **wrong**
 - BTI of java.base requires a local FileSystem instance

```
class FileSystems {  
    static final FileSystem defaultFileSystem;  
    ...  
}
```

GraalVM Native

- BTI is occasionally **desirable** and **wrong**
 - BTI of java.base requires a local FileSystem instance

```
class FileSystems {  
    static final FileSystem defaultFileSystem;  
    ...  
}
```

- Which is set to e.g. a **UnixFileSystem** instance

GraalVM Native

- BTI is occasionally **desirable** and **wrong**
 - BTI of java.base requires a local FileSystem instance

```
class FileSystems {  
    static final FileSystem defaultFileSystem;  
    ...
```

- Which is set to e.g. a UnixFileSystem instance

```
private UnixFileSystemProvider provider;  
private byte[] defaultDirectory;  
private UnixPath rootDirectory;  
...
```

GraalVM Native

- BTI is occasionally **desirable** and **wrong**
 - BTI of java.base requires a local FileSystem instance

```
class FileSystems {  
    static final FileSystem defaultFileSystem;  
    ...
```

- Which is set to e.g. a UnixFileSystem instance

```
private UnixFileSystemProvider provider;  
private byte[] defaultDirectory;  
private UnixPath rootDirectory;  
...
```

- Invalid **private** fields cannot be corrected at runtime!

GraalVM Native Substitutions

GraalVM Native Substitutions

- Graal fixes this by globally rewriting JDK code
 - n.b. to something **equivalent**

GraalVM Native Substitutions

- Graal fixes this by globally rewriting JDK code
 - n.b. to something equivalent
- Relies on Substitution Annotations
 - Recognized by Graal compiler
 - Originally provided for VM transplant
 - i.e. rewriting the internal JDK → JVM API

GraalVM Native Substitutions

- Graal fixes this by globally rewriting JDK code
 - n.b. to something equivalent
- Relies on Substitution Annotations
 - Recognized by Graal compiler
 - Originally provided for VM transplant
 - i.e. rewriting the internal JDK → JVM API
- Very much a sledgehammer approach
 - And needs to be used with care

GraalVM Native Substitutions

```
@TargetClass(classname= "sun.nio.fs.UnixFileSystem")  
final class Target_UnixFileSystem {  
    ...  
}
```

GraalVM Native Substitutions

```
@TargetClass(classname= "sun.nio.fs.UnixFileSystem")
final class Target_UnixFileSystem {
    ...
    @Alias
    Target_UnixFileSystemProvider provider;
    ...
}
```

GraalVM Native Substitutions

```
@TargetClass(classname= "sun.nio.fs.UnixFileSystem")
final class Target_UnixFileSystem {
    ...
    @Alias
    Target_UnixFileSystemProvider provider;
    ...
    @Alias @InjectAccessors(UnixFileSystemAccessors.class)
    private byte[] defaultDirectory;
    @Alias @InjectAccessors(UnixFileSystemAccessors.class)
    private Target_UnixPath rootDirectory;
    ...
}
```

GraalVM Native Substitutions

```
...  
@Inject @RecomputeFieldValue(kind=Kind.Reset)  
byte[] injectedDefaultDirectory;  
@Inject @RecomputeFieldValue(kind=Kind.Reset)  
Target_UnixPath injectedRootDirectory;  
...
```

GraalVM Native Substitutions

```
...  
@Inject @RecomputeFieldValue(kind=Kind.Reset)  
byte[] injectedDefaultDirectory;  
@Inject @RecomputeFieldValue(kind=Kind.Reset)  
Target_UnixPath injectedRootDirectory;  
...  
@Inject @RecomputeFieldValue(kind=Kind.Custom,  
    declClass = NeedsReinitializationProvider.class)  
volatile int needsReinitialization;  
...
```

GraalVM Native Substitutions

```
...  
@Inject @RecomputeFieldValue(kind=Kind.Reset)  
byte[] injectedDefaultDirectory;  
@Inject @RecomputeFieldValue(kind=Kind.Reset)  
Target_UnixPath injectedRootDirectory;  
...  
@Inject @RecomputeFieldValue(kind=Kind.Custom,  
    declClass = NeedsReinitializationProvider.class)  
volatile int needsReinitialization;  
...  
@Alias @TargetElement(name="<init>")  
native void originalConstructor(Target_UnixFileSystemProvider p,  
String dir);  
...
```

GraalVM Native Substitutions

```
class UnixFileSystemAccessors {  
    static void setDefaultDirectory(Target_UnixFileSystem that,  
byte[] value) {  
        that.injectedDefaultDirectory = value;  
    }  
    ...  
}
```


GraalVM Native Substitutions

```
class UnixFileSystemAccessors {  
    static void setDefaultDirectory(Target_UnixFileSystem that,  
byte[] value) {  
        that.injectedDefaultDirectory = value;  
    }  
    static byte[] getDefaultDirectory(Target_UnixFileSystem that) {  
        if (that.needsReinitialization != STATUS_REINITIALIZED) {  
            reinitialize(that);  
        }  
        return that.injectedDefaultDirectory;  
    }  
    ...  
}
```

GraalVM Native Substitutions

```
...  
private static synchronized  
void reinitialize(Target_UnixFileSystem that) {  
    if (that.needsReinitialization !=  
        STATUS_NEEDS_REINITIALIZATION) {  
        return;  
    }  
    that.originalConstructor(that.provider,  
                             System.getProperty("user.dir"));  
    that.needsReinitialization = STATUS_REINITIALIZED;  
}  
...
```

GraalVM Native Substitutions

- **Status quo is inherently risky**
 - Easy to unwittingly introduce bugs
 - Most worrying is security bugs
 - Easy to introduce subtle static vs dynamic disparities

GraalVM Native Substitutions

- **Status quo is inherently risky**
 - Easy to unwittingly introduce bugs
 - Most worrying is security bugs
 - Easy to introduce subtle static vs dynamic disparities
- **But we have a problem**
 - Can we just fix this with Java code changes?
 - Can we fix it with language changes?
 - e.g. (privileged) static reinit paths
 - What to do with legacy code? (JDK, MW and app)

GraalVM Native Startup/Footprint

GraalVM Native Startup/Footprint

- Minimize image code

GraalVM Native Startup/Footprint

- Minimize image code
 - Drop provably uncalled methods
 - Including inline-only methods

GraalVM Native Startup/Footprint

- Minimize image code
 - Drop provably uncalled methods
 - Including inline-only methods
 - Drop BTI-only static init code

GraalVM Native Startup/Footprint

- **Minimize Image Metadata**
 - Drop unreferenced types

GraalVM Native Startup/Footprint

- Minimize Image Metadata
 - Drop unreferenced types
 - Merge Klass into Class<?>

GraalVM Native Startup/Footprint

- **Minimize Image Metadata**
 - Drop unreferenced types
 - Merge Klass into Class<?>
 - GraalVM class DynamicHub
 - Retain only some structural info ...
 - ... size, oop maps, vtables, super/interfaces

GraalVM Native Startup/Footprint

- **Minimize Image Metadata**
 - Drop unreferenced types
 - Merge Klass into Class<?>
 - GraalVM class DynamicHub
 - Retain only some structural info ...
 - ... size, oop maps, vtables, super/interfaces
 - Minimal details of Methods/Fields

GraalVM Native Startup/Footprint

- **Minimize Image Metadata**
 - Drop unreferenced types
 - Merge Klass into Class<?>
 - GraalVM class DynamicHub
 - Retain only some structural info ...
 - ... size, oop maps, vtables, super/interfaces
 - Minimal details of Methods/Fields ...
 - ... e.g. value type, signature, attributes
 - Reflection/Handle info only where needed

GraalVM Native Startup/Footprint

- Minimize Image State

GraalVM Native Startup/Footprint

- Minimize Image State
 - Drop BTI-only static fields

GraalVM Native Startup/Footprint

- **Minimize Image State**
 - Drop BTI-only static fields
 - Drop RO static fields
 - Treat them as global constant
 - Deduplicate repeated primitives & constant objects

GraalVM Native Startup/Footprint

- **Minimize Image State**
 - Drop BTI-only static fields
 - Drop RO static fields
 - Treat them as global constant
 - Deduplicate repeated primitives & constant objects
 - Drop unreferenced instance fields
 - Rare opportunity for smaller data (vs state)

GraalVM Native Startup/Footprint

- Minimize linkage

GraalVM Native Startup/Footprint

- Minimize linkage
 - Code → Code
 - Direct Java ↔ Java calls resolved at image link
 - Direct call for JNI Java -> native impl

GraalVM Native Startup/Footprint

- Minimize linkage
 - Code → Code
 - Direct Java ↔ Java calls resolved at image link
 - Direct call for JNI Java → native impl
 - Metadata ↔ Data ← Code → Metadata
 - Initial Heap objects pinned
 - Constant values pinned
 - Class & class static fields are pinned

GraalVM Native Startup/Footprint

- Minimize linkage
 - Code → Code
 - Direct Java ↔ Java calls resolved at image link
 - Direct call for JNI Java → native impl
 - Metadata ↔ Data ← Code → Metadata
 - Initial Heap objects pinned
 - Constant values pinned
 - Class & class static fields are pinned
 - Very little **load time** mapping/copying & linking

OpenJDK Project Leyden

OpenJDK Project Leyden

- Three Deliverables

OpenJDK Project Leyden

- **Three Deliverables**

- 1 **Extend Java Specifications to Static Java**

- Java Language
- Java Virtual Machine
- JDK Runtime
- Sanctioned Variations and/or Exemptions?

OpenJDK Project Leyden

- **Three Deliverables**

- 1 Extend Java Specifications to Static Java**

- Java Language
- Java Virtual Machine
- JDK Runtime
- Sanctioned Variations and/or Exemptions?

- 2 Extend Java TCK to Static Java**

- Sanctioned Variations and/or Exemptions

OpenJDK Project Leyden

- **Three Deliverables**

- 1 Extend Java Specifications to Static Java**

- Java Language
- Java Virtual Machine
- JDK Runtime
- Sanctioned Variations and/or Exemptions?

- 2 Extend Java TCK to Static Java**

- Sanctioned Variations and/or Exemptions

- 3 Reference Implementation of Static Java?**

Leyden Reference Implementation?

Leyden Reference Implementation

- Reuse (most of) Hotspot JVM in target runtime

Leyden Reference Implementation

- Reuse (most of) Hotspot JVM in target runtime
 - Core runtime (non-dynamic subset)
 - GCs
 - Memory Management Subsystem +
 - Metadata (non-dynamic subset) *
 - Code Cache *

Leyden Reference Implementation

- Reuse (most of) Hotspot JVM in target runtime
 - Core runtime (non-dynamic subset)
 - GCs
 - Memory Management Subsystem +
 - Metadata (non-dynamic subset) *
 - Code Cache *
- Drop unneeded subsystems

Leyden Reference Implementation

- Reuse (most of) Hotspot JVM in target runtime
 - Core runtime (non-dynamic subset)
 - GCs
 - Memory Management Subsystem +
 - Metadata (non-dynamic subset) *
 - Code Cache *
- Drop unneeded subsystems
 - Interpreter & Compilers
 - Class Loading/Bytecode Parsing

Leyden Reference Implementation

- Link generated ELF lib to static libjvm
 - Generated sections at fixed addresses

Leyden Reference Implementation

- **Link generated ELF lib to static libjvm**
 - Generated sections at fixed addresses
- **CodeCache and compiled code**

Leyden Reference Implementation

- **Link generated ELF lib to static libjvm**
 - Generated sections at fixed addresses
- **CodeCache and compiled code**
- **Metadata**
 - Class Model, Symbols, Loader Graph
 - Drop linkage/compile info e.g. CPCache, MethodData

Leyden Reference Implementation

- **Link generated ELF lib to static libjvm**
 - Generated sections at fixed addresses
- **CodeCache and compiled code**
- **Metadata**
 - Class Model, Symbols, Loader Graph
 - Drop linkage/compile info e.g. CPCache, MethodData
- **Initial Heap Region**
 - Pinned contents
 - Extended by GC with dynamic mapped regions

Leyden Experiments

Leyden Experiments

- Reduced VM

Leyden Experiments

- **Reduced VM**
 - **Build libjvm with subsystems excluded**
 - **Compiler**
 - **Interpreter**
 - **Class Loading**

Leyden Experiments

- **Reduced VM**
 - Build libjvm with subsystems excluded
 - Compiler
 - Interpreter
 - Class Loading
 - Test by importing metadata/code from parent JVM
 - Requires substantial cross-linking by hand

Leyden Experiments

- Build Time Closure Analysis

Leyden Experiments

- **Build Time Closure Analysis**
 - **In JVM over CI Interfaces**
 - Less messy than bytecode
 - C1/C2 consume CI model
 - Interfaces mean model is plastic (→ wrappable)

Leyden Experiments

- **Build Time Closure Analysis**
 - **In JVM over CI Interfaces**
 - Less messy than bytecode
 - C1/C2 consume CI model
 - Interfaces mean model is plastic (→ wrappable)
 - **Repurpose C2 to support analysis**
 - Single method graph
 - Inlined method graph
 - Must not deopt

Leyden Experiments

- **Build Time CodeCache generation using C2**
 - **Generate CompiledMethod as per AOT**
 - Or cheat and use existing nmethod
 - Which can be laid out as as per AOT methods

Leyden Experiments

- **Build Time CodeCache generation using C2**
 - **Generate CompiledMethod as per AOT**
 - Or cheat and use existing nmethod
 - Which can be laid out as as per AOT methods
 - **Must not deopt cold paths or speculate**
 - n.b. closed world → speculation becomes determinate

Leyden Experiments

- **Build Time CodeCache generation using C2**
 - **Generate CompiledMethod as per AOT**
 - Or cheat and use existing nmethod
 - Which can be laid out as as per AOT methods
 - **Must not deopt cold paths or speculate**
 - n.b. closed world → speculation becomes determinate
 - **Resolve all calls at compile time**
 - C2 does late resolution of call sites
 - **Ideally JVM only links to CodeCache static fields**

Leyden Experiments

- **Build Time Metadata generation**
 - Similar to current CDS but as ELF lib
 - .metadata section pinned at fixed address
 - Cannot miss out 'difficult' classes
 - Save whole tree from `ClassLoaderDataGraph::head`

Leyden Experiments

- **Build Time Metadata generation**
 - Similar to current CDS but as ELF lib
 - .metadata section pinned at fixed address
 - Cannot miss out 'difficult' classes
 - Save whole tree from `ClassLoaderDataGraph::head`
 - **Generate ELF relocs for all references**
 - Internal links, initial heap data, method/code pointers

Leyden Experiments

- **Build Time Metadata generation**
 - Similar to current CDS but as ELF lib
 - .metadata section pinned at fixed address
 - Cannot miss out 'difficult' classes
 - Save whole tree from `ClassLoaderDataGraph::head`
 - Generate ELF relocs for all references
 - Internal links, initial heap data, method/code pointers
 - Ideally JVM only links in Metadata static fields
 - Plus minor memory region init and validation

Leyden Experiments

- **Build Time Initial Heap generation**
 - Also similar to current CDS but as ELF lib
 - .heap_data section pinned at fixed address
 - RO subsection for constants
 - RW subsection for objects (including mirrors)

Leyden Experiments

- **Build Time Initial Heap generation**
 - Also similar to current CDS but as ELF lib
 - .heap_data section pinned at fixed address
 - RO subsection for constants
 - RW subsection for objects (including mirrors)
 - **Generate ELF relocs for all references**
 - Internal links, metadata (Klass) pointers

Leyden Experiments

- **Build Time Initial Heap generation**
 - Also similar to current CDS but as ELF lib
 - .heap_data section pinned at fixed address
 - RO subsection for constants
 - RW subsection for objects (including mirrors)
 - Generate ELF relocs for all references
 - Internal links, metadata (Klass) pointers
 - GC must include as a heap sub-region
 - Avoids copy and link reloc as per current CDS

Leyden Experiments Summary

- **Reduced VM**
 - Attempted – very difficult to decouple unused code
- **Generated CodeCache**
 - Attempted – Saved and reloaded ELF lib
 - Metadata/object data still hand linked
- **CI Closure Generation**
 - Started – mostly been working on closure and init analysis
- **Metadata/Initial Heap**
 - Still to do

Thank You!

Questions?