

# Static Java, GraalVM Native and OpenJDK

Andrew Dinn  
Red Hat Java Team  
October 2021

# Static Java?

# Static Java?

- **A New Deployment Model?**
  - Pre-compiled, self-contained executable like C++/Go
    - No class files → no interpreter/JIT, less metadata ...
- **A Container-Friendly Java?**
  - Smaller footprint/faster startup
    - Not automatic
- **A Reduced Java?**
  - Dynamic (runtime) → Static (buildtime)
    - Java? or just the JVM/JDK?

# Dynamic Java

# How Can Java Be Static?

- **Program code evolves during execution**
  - Class loading, Service providers, Class generators
  - Loader Delegation, Module Layers
- **Execution model supports this evolution**
  - Reflection, Method/VarHandles
    - With ClassLoader → indirect load and link

# How Can Java Be Static?

- Program state evolves too
  - { Load → Static Init → Interpret → (Re-)Compile } +
    - n.b. Compile happens after Static Init/Execute
- Static Init may recursively embed this pattern
  - Static Init → { Load → ... } +
- Load also drives recompile
  - Load → Deopt(dependent) → Interpret → ReCompile

# How Can Java Be Static?

- **VM configured at runtime**
  - GC, JITs, Compressed Oops, HW Accelerators, etc
    - Configurable on command line
    - Defaulted according to runtime env
- **Tools**
  - Agents, JDWP, JFR, JMX, jcmd, etc
    - Can be installed & configured at runtime
    - All rely on dynamic capabilities

# How Can We Make Java Static?



# How To Make Java Static?

- **Close the world of the app**
  - All code identified and compiled ahead of time
    - Trace class, method and field references from App main
    - Try not to include the kitchen sink
- **Configure indirectly accessed code?**
  - Reflection & method/varhandle targets
  - Service implementations
- **Configure classpath resources?**

# How To Make Java Static?

- Close the world of the target runtime
  - But do it at build time
- Can we emulate the dynamic runtime?
  - App-defined loaders in build env?
    - Or just the classpath & modulepath?
  - Which services get loaded?
    - e.g. what will the target LANG setting be

# How To Make Java Static?

- No Load → No static init at load?
  - So when do we run class static init?

# How To Make Java Static?

- No Load → No static init at load?

## 1 Emulate dynamic JVM?

- Run init at point(s) where load *would* have happened?
  - Including init that precedes App main
  - Check at every new or static access/call
    - Or at mention of Foo.class or `forName()` call
  - Init only once (however we get there)
    - Analysis can remove many redundant attempts

# How To Make Java Static?

- No Load → No static init at load
- ## 2 Run all inits in one go before App main?
- But still in a 'correct' order?
  - Might this change the init outcome relative to 1?  
`A::init(); A.STATIC_VALUE = X; B::init()`
  - Maybe some hybrid of 1 and 2 could avoid that
    - Needs a very careful analysis

# How To Make Java Static?

- No Load → No static init at load
- ## 3 Pre-compute static fields at build time?
- As appropriate to the target env
    - See above ...
      - ... but with clairvoyance requirement
    - Actually a hybrid approach must be adopted ...
      - ... and sometimes you need to cheat?

# How To Make Java Static?

- **Build time init is fine for many fields**
  - Final primitive constants
  - Immutable objects (e.g. Strings)
  - Final mutable objects?
- **BTI even opens up new compiler optimizations**
  - Non-final values that are read-only in app closure
    - Effectively final

# How To Make Java Static?

- But some values really must be runtime initied
  - Thread, FileStream
  - Random, Credentials, etc
  - Values computed from native calls/JVM/OS env
  - Values derived from any/all of the above



# How To Make Java Static?

- It matters how values are computed/consumed
  - `public static final int UnsafeConstants.PAGE_SIZE = 0;`
    - JVM bootstrap injects real value
    - Used/exposed via other APIs? e.g. `ByteBuffer`?
  - Default/Alternative Locales
    - Choice drives init of text processing subsystem
  - Default `FileSystemProvider`
    - Associated `FileSystem` caches local root/working dir

# How To Make Java Static?

- **JVM Configuration**
  - Configure compilation for target JVM environment
    - GC, Compressed Oops, HW Accelerators, etc
- **Tools Configuration**
  - Convert some tools to operate at build time?
    - Agents, JFR
      - Graal: JFR with a runtime on/off switch
  - Implement complete replacements for others?
    - Graal: JDWP → DWARF, JMX/jcmd → ???

# GraalVM Native

# GraalVM Native

- **Closure Analysis**
  - Derives full reference closure
  - Able to prune unused code and (instance/static) state
    - More or less aggressively (high build time cost)
- **Class/Module Loading**
  - Standard ClassLoaders exist but don't load
    - Return existing class or throw CNFE
  - Queryable Module base layer info recently added

# GraalVM Native

- **Reflection**
  - Analysis infers Method/Field target where possible
  - Fall back to user configuration
- **Method Handles**
  - Handle chains must end at `DirectHandle`
    - Lambdas still work
- **Service & Resource Loading**
  - Analysis infers target where possible
  - Fall back to user configuration

# GraalVM Native

- **Special cased JDK Service Providers**
  - **Locale Support**
    - User configures default locale at build
      - Must also pre-specify all other required locales
  - **Default File System**
    - Bootstrap requires default Provider & FileSystem
    - “Optimized” by creating at build-time
      - Requires patching JDK class

# GraalVM Native

- Almost all of JDK is build-time init'd
  - Some of it is always runtime init'd
  - Some of it gets re-init'd
    - Which is a certainly questionable?
- Apps can be configured for BTI
  - Almost all frameworks try to do this
- Static init analysis rejects invalid cases
  - Values that must be RTI
  - Values that derive from RTI values

# GraalVM Native

- **BTI is critical to small JDK footprint/fast startup**
  - Only static state referenced by app is retained
    - Intermediate computed values squeezed out
      - As is their type info
  - Only RTI static init code is compiled/executed
    - Especially JDK default locale and file system setup
    - This is a big slice of dynamic JVM startup
  - n.b. the same story applies for app code



# GraalVM Native

- Example of REINIT

```
class JDKInitializationFeature {  
    ...  
    RuntimeClassInitializationSupport rci = ...;  
    rci.initializeAtBuildTime("java.io",  
                             "JDK core init'd at build time");  
    rci.initializeAtBuildTime("java.lang",  
                             "JDK core init'd at build time");  
    ...  
    rci.rerunInitialization("sun.nio.ch.Epoll",  
                           "Calls Epoll.eventSize(), ...");  
}
```

# GraalVM Native

- Example of REINIT

```
class Epoll {  
    ...  
    private static final  
    int SIZEOF_EPOLLEVENT    = eventSize();  
    ...  
    // opcodes  
    static final int EPOLL_CTL_ADD = 1;  
    ...  
    private static native int eventSize();  
}
```

# GraalVM Native

- BTI is occasionally **desirable** and **wrong**
  - BTI of java.base requires a local FileSystem instance

```
class FileSystems {  
    static final FileSystem defaultFileSystem;  
    ...  
}
```

- Which is set to e.g. a UnixFileSystem instance

```
private UnixFileSystemProvider provider;  
private byte[] defaultDirectory;  
private UnixPath rootDirectory;  
...
```

- Invalid **private** fields cannot be corrected at runtime!

# GraalVM Native Substitutions

# GraalVM Native Substitutions

- Graal fixes this by globally rewriting JDK code
  - n.b. to something equivalent
- Relies on Substitution Annotations
  - Recognized by Graal compiler
  - Originally provided for VM transplant
    - i.e. rewriting the internal JDK → JVM API
- Very much a sledgehammer approach
  - And needs to be used with care

# GraalVM Native Substitutions

```
@TargetClass(classname= "sun.nio.fs.UnixFileSystem")
final class Target_UnixFileSystem {
    ...
    @Alias
    Target_UnixFileSystemProvider provider;
    ...
    @Alias @InjectAccessors(UnixFileSystemAccessors.class)
    private byte[] defaultDirectory;
    @Alias @InjectAccessors(UnixFileSystemAccessors.class)
    private Target_UnixPath rootDirectory;
    ...
}
```

# GraalVM Native Substitutions

```
...
@Inject @RecomputeFieldValue(kind=Kind.Reset)
byte[] injectedDefaultDirectory;
@Inject @RecomputeFieldValue(kind=Kind.Reset)
Target_UnixPath injectedRootDirectory;
...
@Inject @RecomputeFieldValue(kind=Kind.Custom,
    declClass = NeedsReinitializationProvider.class)
volatile int needsReinitialization;
...
@Alias @TargetElement(name="<init>")
native void originalConstructor(Target_UnixFileSystemProvider p,
String dir);
...
```

# GraalVM Native Substitutions

```
class UnixFileSystemAccessors {  
    static void setDefaultDirectory(Target_UnixFileSystem that,  
byte[] value) {  
        that.injectedDefaultDirectory = value;  
    }  
    static byte[] getDefaultDirectory(Target_UnixFileSystem that) {  
        if (that.needsReinitialization != STATUS_REINITIALIZED) {  
            reinitialize(that);  
        }  
        return that.injectedDefaultDirectory;  
    }  
    ...  
}
```



# GraalVM Native Substitutions

```
...  
private static synchronized  
void reinitialize(Target_UnixFileSystem that) {  
    if (that.needsReinitialization !=  
        STATUS_NEEDS_REINITIALIZATION) {  
        return;  
    }  
    that.originalConstructor(that.provider,  
                             System.getProperty("user.dir"));  
    that.needsReinitialization = STATUS_REINITIALIZED;  
}  
...
```

# GraalVM Native Substitutions

- **Status quo is inherently risky**
  - Easy to unwittingly introduce bugs
    - Most worrying is security bugs
  - Easy to introduce subtle static vs dynamic disparities
- **But we have a problem**
  - Can we just fix this with Java code changes?
  - Can we fix it with language changes?
    - e.g. (privileged) static reinit paths
  - What to do with legacy code? (JDK, MW and app)

# GraalVM Native Startup/Footprint

# GraalVM Native Startup/Footprint

- Minimize image code
  - Drop provably uncalled methods
    - Including inline-only methods
  - Drop BTI-only static init code

# GraalVM Native Startup/Footprint

- **Minimize Image Metadata**
  - Drop unreferenced types
  - Merge Klass into Class<?>
    - GraalVM class DynamicHub
    - Retain only some structural info ...
      - ... size, oop maps, vtables, super/interfaces
  - Minimal details of Methods/Fields ...
    - ... e.g. value type, signature, attributes
    - Reflection/Handle info only where needed

# GraalVM Native Startup/Footprint

- **Minimize Image State**
  - Drop BTI-only static fields
  - Drop RO static fields
    - Treat them as global constant
    - Deduplicate repeated primitives & constant objects
  - Drop unreferenced instance fields
    - Rare opportunity for smaller data (vs state)

# GraalVM Native Startup/Footprint

- Minimize linkage
  - Code → Code
    - Direct Java ↔ Java calls resolved at image link
    - Direct call for JNI Java → native impl
  - Metadata ↔ Data ← Code → Metadata
    - Initial Heap objects pinned
    - Constant values pinned
    - Class & class static fields are pinned
  - Very little load time mapping/copying & linking

# OpenJDK Project Leyden



# OpenJDK Project Leyden

- **Three Deliverables**

- 1 Extend Java Specifications to Static Java**

- Java Language
- Java Virtual Machine
- JDK Runtime
- Sanctioned Variations and/or Exemptions?

- 2 Extend Java TCK to Static Java**

- Sanctioned Variations and/or Exemptions

- 3 Reference Implementation of Static Java?**

# Leyden Reference Implementation?

# Leyden Reference Implementation

- Reuse (most of) Hotspot JVM in target runtime
  - Core runtime (non-dynamic subset)
  - GCs
  - Memory Management Subsystem +
  - Metadata (non-dynamic subset) \*
  - Code Cache \*
- Drop unneeded subsystems
  - Interpreter & Compilers
  - Class Loading/Bytecode Parsing

# Leyden Reference Implementation

- **Link generated ELF lib to static libjvm**
  - Generated sections at fixed addresses
- **CodeCache and compiled code**
- **Metadata**
  - Class Model, Symbols, Loader Graph
  - Drop linkage/compile info e.g. CPCache, MethodData
- **Initial Heap Region**
  - Pinned contents
  - Extended by GC with dynamic mapped regions

# Leyden Experiments

# Leyden Experiments

- **Reduced VM**
  - Build libjvm with subsystems excluded
    - Compiler
    - Interpreter
    - Class Loading
  - Test by importing metadata/code from parent JVM
    - Requires substantial cross-linking by hand

# Leyden Experiments

- **Build Time Closure Analysis**
  - **In JVM over CI Interfaces**
    - Less messy than bytecode
    - C1/C2 consume CI model
    - Interfaces mean model is plastic (→ wrappable)
  - **Repurpose C2 to support analysis**
    - Single method graph
    - Inlined method graph
    - Must not deopt

# Leyden Experiments

- **Build Time CodeCache generation using C2**
  - **Generate CompiledMethod as per AOT**
    - Or cheat and use existing nmethod
    - Which can be laid out as as per AOT methods
  - **Must not deopt cold paths or speculate**
    - n.b. closed world → speculation becomes determinate
  - **Resolve all calls at compile time**
    - C2 does late resolution of call sites
  - **Ideally JVM only links to CodeCache static fields**



# Leyden Experiments

- **Build Time Metadata generation**
  - Similar to current CDS but as ELF lib
    - .metadata section pinned at fixed address
    - Cannot miss out 'difficult' classes
    - Save whole tree from ClassLoaderDataGraph::head
  - **Generate ELF relocs for all references**
    - Internal links, initial heap data, method/code pointers
  - **Ideally JVM only links in Metadata static fields**
    - Plus minor memory region init and validation

# Leyden Experiments

- **Build Time Initial Heap generation**
  - Also similar to current CDS but as ELF lib
    - .heap\_data section pinned at fixed address
    - RO subsection for constants
    - RW subsection for objects (including mirrors)
  - **Generate ELF relocs for all references**
    - Internal links, metadata (Klass) pointers
  - **GC must include as a heap sub-region**
    - Avoids copy and link reloc as per current CDS

# Leyden Experiments Summary

- **Reduced VM**
  - Attempted – very difficult to decouple unused code
- **Generated CodeCache**
  - Attempted – Saved and reloaded ELF lib
    - Metadata/object data still hand linked
- **CI Closure Generation**
  - Started – mostly been working on closure and init analysis
- **Metadata/Initial Heap**
  - Still to do

**Thank You!**

**Questions?**